Carbon OCR

The fast OCR engine in C++ and CUDA

Measured, Not Claimed

Three pages · One laptop CPU · Every picture drawn from the engine's own output

---

FULLERENE AT WORK

Three pages, read by fullerene-3M

A page of an annual report, a page of two tables, and a rough photo of the first, upside down and tilted: how the engine turns and levels a page, finds its lines, reads them, orders and labels its blocks, places its tables' cells and writes HTML. Every number is measured, every fault is listed with where it comes from, and the scripts make it all again.



Karan Vora, Founder · Carbon Research

Contents

—	At a glance	3
I	**PART I · PAGES, READ**	5
01	The run	6
02	Turned and levelled	7
03	What the detector sees	8
04	Lines, blocks and reading order	9
05	Tables	11
06	The HTML it writes	12
II	**PART II · MEASURED**	13
07	The numbers	14
08	What it gets wrong	15
09	Three pages, and 2,937	16
10	Run it yourself	16

At a glance

Three pages, read end to end on a laptop's CPU. A page of an annual report, a page of two tables, and a rough photo of the first. fullerene-3M turned and levelled the photo, found every line, read the words, labelled and ordered the blocks, placed every table cell and wrote HTML. The pages are invented, so their text is known word for word and none of it was trained on.

ORIENT	the photo scored 180° and was turned, then levelled by 2.45° of the 2.5° it was off, with no model in that step · the clean pages left alone · 27 ms on the photo
DETECT	`det` at 1024 px · 47 lines on the report page and 85 on the page of tables, every printed line · the figure found as a picture · 44 ms on one core
READ	`rec-latin` · 253 of 255 words right on the page and on its photo, 168 of 173 on the tables · 98 ms for the report page's 47 lines on one core
ORDER	`order` · 18 blocks in reading order with their labels · both tables read as grids, all 73 cells in their row and column, 59 exact
WRITE	HTML: headers, headings, paragraphs, lists, figures, tables · or text, markdown, JSON, hOCR, ALTO, PAGE XML and searchable PDF

253 / 255
words right on the report page and on its photo; the two misses are one footnote mark

73 / 73
table cells in their row and column, on a page of two tables; 59 character for character

152 ms
the report page on one CPU core, every model; 60 ms on six

0.05°
left of the photo's tilt once levelled, after a half turn the right way up

7 of 11
of the faults found can be fixed in the engine, with no new training

WHAT IT GETS WRONG

The words are almost all right; the page's structure is where the faults are. Headings joined to the paragraph under them, the title to its subtitle, the footnote to the last paragraph; the table placed after the footers; a column rule read as a "|" beside a cell's number; a caption written as a table's last row; a list per item. Section 8 lists each with where it comes from and what fixes it.

HOW IT IS MEASURED

Words are counted strictly, as runs between spaces, order aside: a word with anything stuck to it is wrong. Cells are compared at their row and column. Times are medians of seven runs after a warm-up, by the engine's own clock; the laptop's clock moves them, 135 to 164 ms for the same page across the day.



PART I

Pages, read

Three pages through the engine: turned and levelled, their lines found, their blocks ordered, their tables read cell by cell, and the HTML it writes.

01 The run

02 Turned and levelled

03 What the detector sees

04 Lines, blocks and reading order

05 Tables

06 The HTML it writes

01 The run

the models	`fullerene-3M` from Hugging Face: `orient`, `det`, `rec-latin` and `order`, 2.77 M parameters, under the Carbon-OCR OpenRAIL-AMS license; the FP32 files on ONNX Runtime, the engine's choice on a CPU for a set without int8 files
the engine	Carbon-OCR as of 3 November 2026, through its Python module: `carbon_ocr.Engine(set, workers=1)` and `workers=6`, `ocr(pages, html=True)`
the machine	a laptop: Intel Core i7-13620H (6 performance cores, 4 efficient), Windows 11, CPU only
the inputs	`page.png` and `tables.png`, the two HTML pages printed by headless Chrome at 1700 × 2200 (a Letter page at 200 dpi); `photo.png`, 1796 × 2274, made from `page.png` (below)
the times	medians of seven runs after a warm-up. The laptop's clock follows its power limit and its heat, so the same page took 135 to 164 ms on one core across the day's runs; the numbers here are the run that drew the pictures

The report page (demo/page.html) is a page of a water district's annual report in two columns: a running header and footer, a title and a subtitle, four headings, five paragraphs, a list of four items, a three-column table, a figure with its caption, and a footnote. The page of tables (demo/tables.html) is the same district's quarterly figures: a title, two short paragraphs, a heading, a table of six columns and eight rows with a bold total, a table of five columns and five rows, and their captions. The photo is the report page turned 182.5°, so upside down and 2.5° off, on paper tinted under a light that brightens to the right, with noise, a slight blur and JPEG at quality 60.

02 Turned and levelled

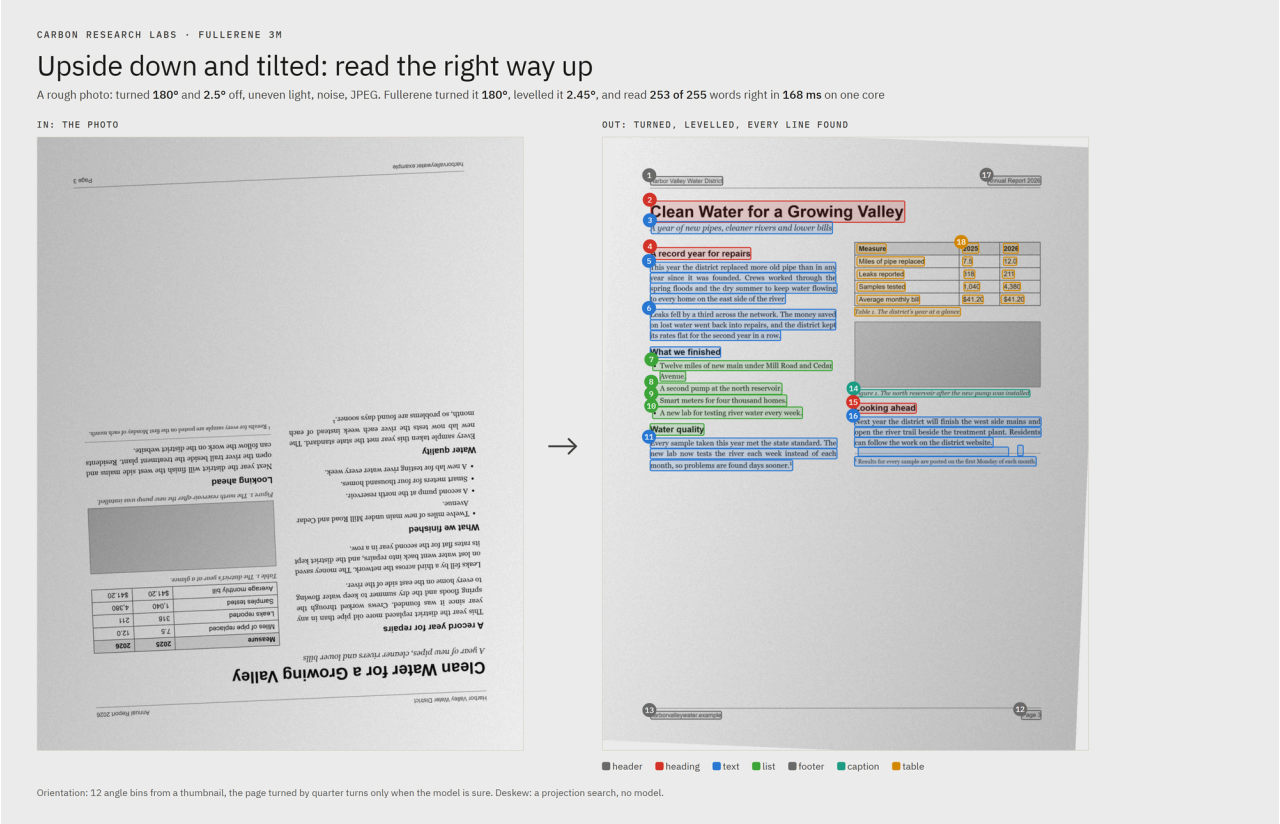


Figure 1. The photo, upside down and 2.5 degrees off, and the page as the engine read it: turned 180 degrees, levelled by 2.45 degrees, every line boxed in its block's colour and numbered in reading order.

Before it looks for a line, the engine scales the page to the detector's size once, and `orient` reads a 512 × 512 thumbnail of that copy: 12 angle bins and a vertical flag. The engine turns the page only by a quarter turn, and only when the winning bin is 90°, 180° or 270° and its probability passes the threshold stamped in `orient.onnx` (0.65). The photo scored 180° and was turned; the clean pages scored upright and were left alone.

What is left of a tilt is found without a model: a projection search over small angles on the detector-scale copy, kept when the level angle clearly wins. It found 2.45° of the 2.5° the photo was off, and levelled the page and its copy by it. This whole step (the scaling, the thumbnail, the model, the turn and the levelling) is 11 ms on a clean page and 27 ms on the photo, where the whole page is turned and then rotated.

The photo's reading is as good as the clean page's: the same 253 of 255 words, 48 lines against 47, and a mean line confidence of 0.986 against 0.992.

03 What the detector sees

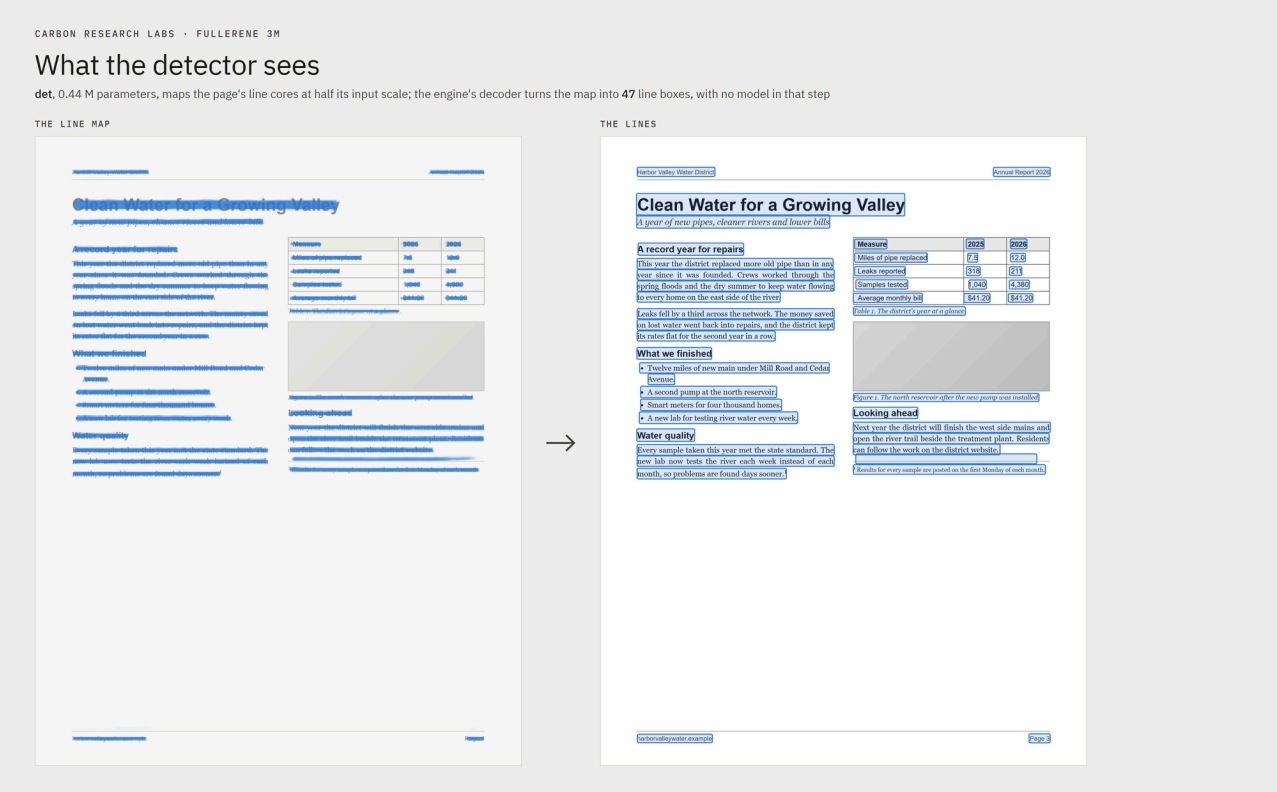


Figure 2. The detector's line-core map over the report page, and the 47 line boxes the engine's decoder made of it.

`det` reads the page at 1024 px on its long side (791 × 1024 here, padded to 800 × 1024) and writes four maps at half that scale: each line's core band (the middle 40% of its height, shortened at both ends), the distances from the band to the line's top and bottom, and pictures. The left half of the figure is the core map over the page. The engine's decoder thresholds it, takes the connected bands, and sets each line's top and bottom from the distance maps, in C++ with no model: 0.6 ms for the page.

It found 47 lines on the report page, every printed line and one more: the footnote's rule, a thin line the recognizer then reads as underscores (confidence 0.69, the page's lowest). The figure is found as a picture, one box on the grey panel, so no line is looked for inside it. On the page of tables it found 85 lines, a line for every cell. Detection is 44 ms on one core and 15 ms on six, where the page's detection runs on the cores the other workers leave idle.

04 Lines, blocks and reading order

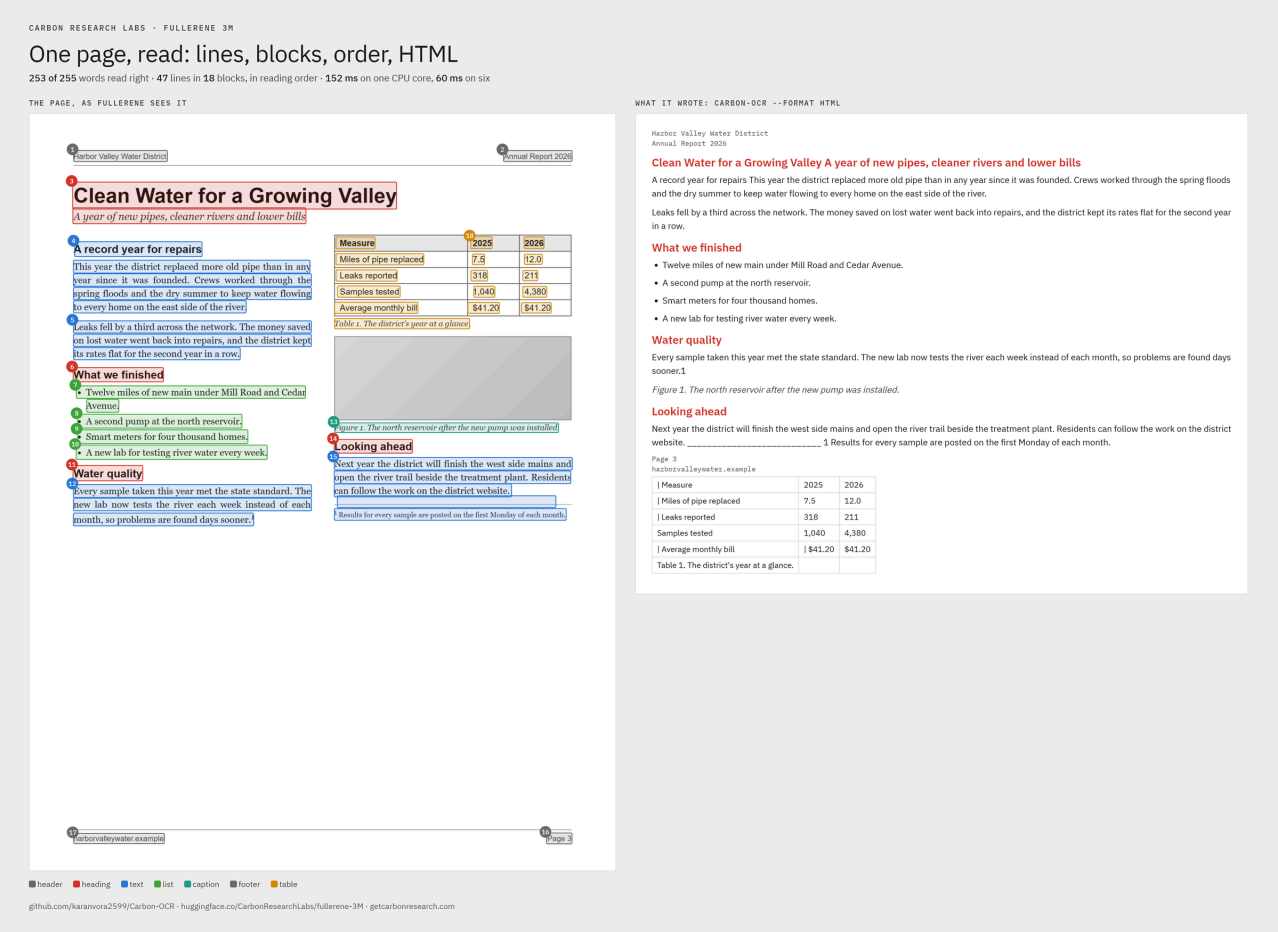


Figure 3. The report page with every line boxed in its block's colour and every block numbered in reading order, beside the HTML the engine wrote for it.

`order` reads 16 numbers for each line (its box, its size and shape against the page's, and a few things about its text: its length, its share of digits, a list marker, a trailing hyphen, a sentence end, a capital start) and gives, for every pair of lines, how likely one follows the other and whether they share a block; for each line one of 13 block labels; and the table structure, as which lines share a row, a column, a cell. The engine decodes it in C++: a chain through the lines from the most likely first one, repaired by moving a line wherever the chain's probability rises most; blocks as the runs of that chain whose neighbours share a block; and each block's label as the mean of its lines'. When the chain is not sure enough the XY-cut, a rule on the boxes alone, orders the page instead; on all three pages the model's order was kept.

The report page's 18 blocks, in the order the engine gave them:

#	LABEL	WHAT IT HOLDS	BLOCK CONFIDENCE
1, 2	page header	the district's name; "Annual Report 2026"	1.00, 1.00
3	heading	the title and the subtitle together	0.62

#	LABEL	WHAT IT HOLDS	BLOCK CONFIDENCE
4	text	the heading "A record year for repairs" and the paragraph under it	0.70
5	text	the second paragraph	0.94
6	heading	"What we finished"	0.98
7 to 10	list	the four items, one block each	0.99 to 1.00
11	heading	"Water quality"	0.97
12	text	its paragraph	0.95
13	caption	the figure's caption	0.95
14	heading	"Looking ahead"	0.68
15	text	its paragraph, then the footnote's rule and the footnote	0.95
16, 17	page footer	"Page 3"; the site	1.00, 0.99
18	table	the table's 15 cells and its caption	0.98

The columns are read in order, left then right, and the figure's caption, the headings and the list are where a reader would put them. The faults are in the same table: the blocks that joined two kinds of line (3, 4 and 15) and the table's place at the end (section 7). The blocks that joined lines wrongly are also the ones the model was least sure of.

05 Tables

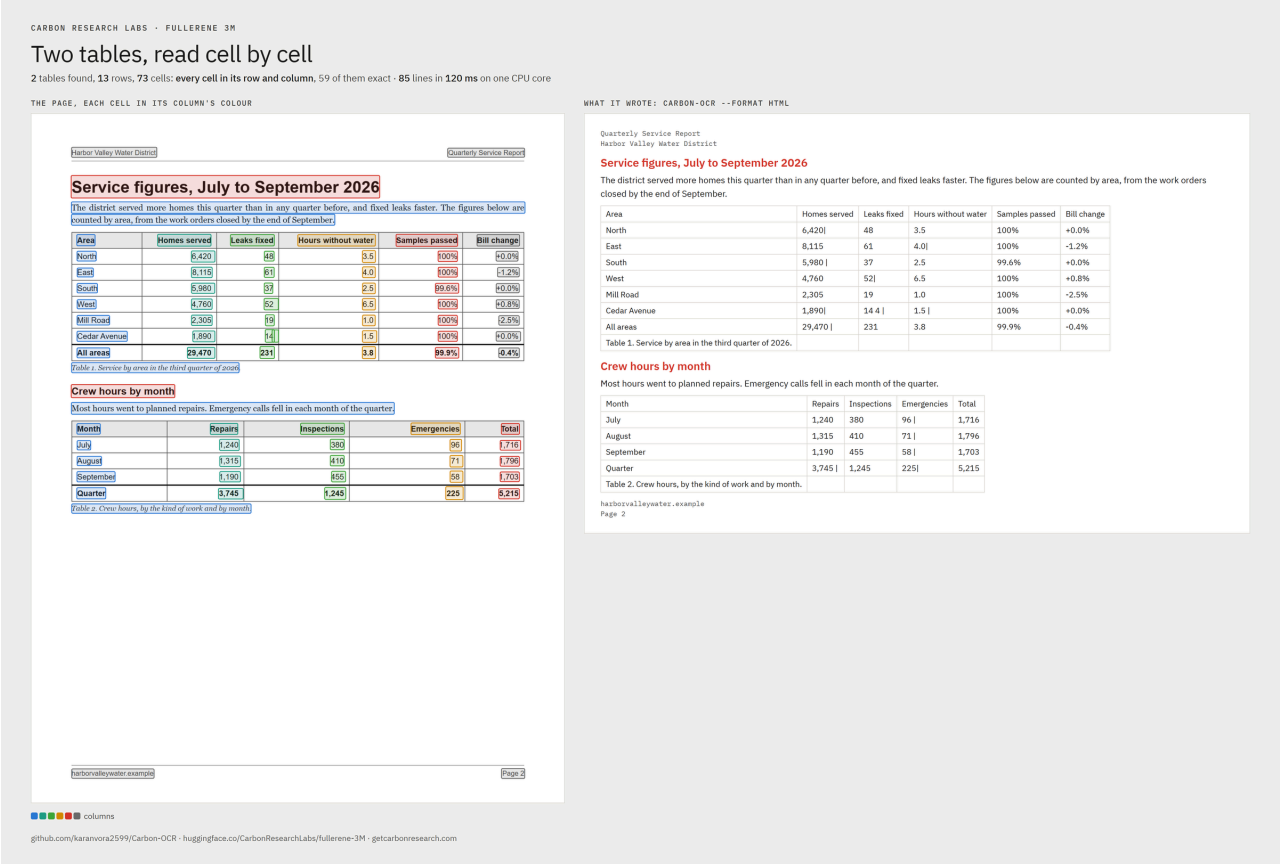


Figure 4. The page of two tables with every cell boxed in its column's colour, beside the HTML tables the engine wrote: every cell in its row and column.

Each table is a block whose lines `order` groups into rows, columns and cells; the engine reads the groups as a grid and writes it as an HTML table, row by row. On the page of tables both tables were found as tables, and all 73 cells landed in their own row and column, the header rows and the bold totals included:

	TABLE 1	TABLE 2	BOTH
rows, columns	8 × 6	5 × 5	13 rows
cells	48	25	73
in their row and column	48	25	**73**
exact, character for character	39	20	**59**
right but for the column rule's `|` beside the number, or a stray space	8	5	13
right, with a stray fragment beside it	1 ("14 4 |" for 14)	0	1
rows written	9 (the caption as a ninth)	6 (the same)	

On the report page the three-column table is read the same way: its 15 cells in place, five of them with a `|`.

06 The HTML it writes

The report page as `ocr(pages, html=True)` (or `carbon-ocr ocr --format html`) wrote it, its long lines wrapped here:

```
<header><p>Harbor Valley Water District</p></header>
<header><p>Annual Report 2026</p></header>
<h2>Clean Water for a Growing Valley A year of new pipes, cleaner rivers and lower
  bills</h2>
<p>A record year for repairs This year the district replaced more old pipe than in any year
  since it was founded. Crews worked through the spring floods and the dry summer to keep
  water flowing to every home on the east side of the river.</p>
<p>Leaks fell by a third across the network. The money saved on lost water went back into
  repairs, and the district kept its rates flat for the second year in a row.</p>
<h2>What we finished</h2>
<ul><li>Twelve miles of new main under Mill Road and Cedar Avenue.</li></ul>
<ul><li>A second pump at the north reservoir.</li></ul>
<ul><li>Smart meters for four thousand homes.</li></ul>
<ul><li>A new lab for testing river water every week.</li></ul>
<h2>Water quality</h2>
<p>Every sample taken this year met the state standard. The new lab now tests the river each
  week instead of each month, so problems are found days sooner.1</p>
<figure><figcaption>Figure 1. The north reservoir after the new pump was
  installed.</figcaption></figure>
<h2>Looking ahead</h2>
<p>Next year the district will finish the west side mains and open the river trail beside
  the treatment plant. Residents can follow the work on the district website.
  _____ 1 Results for every sample are posted on the first Monday of
  each month.</p>
<footer><p>Page 3</p></footer>
<footer><p>harborvalleywater.example</p></footer>
<table>
<tr><td>| Measure</td><td>2025</td><td>2026</td></tr>
<tr><td>| Miles of pipe replaced</td><td>7.5</td><td>12.0</td></tr>
<tr><td>| Leaks reported</td><td>318</td><td>211</td></tr>
<tr><td>Samples tested</td><td>1,040</td><td>4,380</td></tr>
<tr><td>| Average monthly bill</td><td>| $41.20</td><td>$41.20</td></tr>
<tr><td>Table 1. The district's year at a glance.</td><td></td><td></td></tr>
</table>
```

Each block becomes one element by its label: page headers and footers `<header>` and `<footer>`, headings `<h2>`, text `<p>`, a list `<ul>`, a caption a `<figure>` with its `<figcaption>`, a table a `<table>`. A block's lines are joined into one paragraph, and a word broken by a hyphen at a line's end is joined again. The bullets are dropped from list items, where the list's own markers stand for them. The same page as plain text, markdown, JSON (with every line's box and confidence), hOCR, ALTO, PAGE XML or a searchable PDF is one option away.



PART II

Measured

The numbers, every fault and where it comes from, the benchmark behind them, and how to run it all again.

07 The numbers

08 What it gets wrong

09 Three pages, and 2,937

10 Run it yourself

07 The numbers

	REPORT PAGE	ITS PHOTO	PAGE OF TABLES
words in the source	255	255	173
words read	264	262	182
words read right	**253 (99.2%)**	**253 (99.2%)**	168 (97.1%)
of the words read, right	95.8%	96.6%	92.3%
lines	47	48	85
blocks	18	18	10
mean line confidence	0.992	0.986	0.988
turned, levelled	0°, 0°	180°, 2.45°	0°, 0°

Words are counted as runs of characters between spaces, as a multiset, order aside, and strictly: a word with anything stuck to it is wrong. The misses are what is stuck to words. On the report page both are one footnote mark, read as a plain 1 into the word before it ("sooner.1" for "sooner." and its mark); on the page of tables, the five misses are numbers with a column rule's `|` against them ("6,420|"). The extra words are marks with no word in the source: the list's four bullets, the table rules read as `|` (five on the report page, two on the photo, eight on the page of tables), the footnote's rule as underscores (and a second short run of them on the photo), and one stray "4".

Time, in milliseconds, each stage by the engine's own clock and the whole call from Python:

STAGE	REPORT, ONE CORE	REPORT, SIX	PHOTO, ONE CORE	PHOTO, SIX	TABLES, ONE CORE	TABLES, SIX
scale, orient, turn, level	11.5	11.5	27.1	26.9	10.5	11.5
detect	44.1	15.0	44.2	15.7	38.4	14.7
decode the maps	0.6	0.6	0.8	0.9	0.5	0.5
recognize	97.7	31.3	91.0	33.1	69.2	38.1
the page	**151.8**	**60.5**	**168.3**	**77.3**	**120.2**	**67.3**

Each row is a median of its own, so the rows do not add up to the page exactly; what follows recognition (the order model, 22 MMAC a page, and the output) is inside that rounding. One page alone uses all six workers: its detection runs on the idle cores, and its lines are split among them, so the report page reads 2.5 times as fast on six as on one. Recognition is the most work: 47 lines, about two-thirds of the report page's time on one core. The page of tables has more lines (85) but shorter ones, so it reads faster.

08 What it gets wrong

FAULT	WHERE	WHERE IT COMES FROM	WHAT FIXES IT
the title and the subtitle are one heading	report page	`order` put the two lines in one block, and the block's label is their mean (block confidence 0.62)	in the engine: split a block where the line height or the label changes sharply between two lines; in training, more titles with subtitles
a heading inside the paragraph under it	report page: "A record year for repairs" (0.70); photo: "What we finished" in a paragraph (0.83), "Water quality" in the last list item (0.53)	the same: one block for two kinds of line, the heading outvoted by the lines under it	the same split; the model's low block confidence marks every one of these blocks
the footnote inside the last paragraph, its rule as underscores	report page and photo	`det` finds the rule as a line, and `order` puts the rule and the footnote in the paragraph's block	in the engine: drop a line that is only underscores or dashes before ordering; in training, footnotes under their rule
the table comes last, after the footers	report page and photo	the order's chain puts the table's lines at the end	in training: tables mid-column, the layouts the XY-cut also gets wrong
a `|` in a cell, beside its text	5 cells on the report page, 2 on the photo, 13 on the page of tables	a cell's crop takes in the column rule beside it, and `rec-latin` reads the rule as `|`	in the engine: trim a `|` at a cell's edge, or crop cells inside their rules
a stray fragment in a cell	one cell of 73 ("14 4 |")	a second, partial line found at the cell's rule	the same trimming, and dropping a cell's line that overlaps another
a table's caption written as its last row	both pages	`order` puts the caption's line in the table's block	in the engine: a line under a table that spans its columns and fills only the first is its caption
each list item a list of its own	report page and photo	each item is its own list block, and the writer makes a `<ul>` of each	in the engine: the writer joins adjacent list blocks into one `<ul>`
eight picture regions that are not pictures, one of them the right half of the page	photo	`det`'s picture map takes the paper's edge and the light falling off it for a picture	in training, photographed pages, their edges and backgrounds; the lines inside are still found and read
on the photo the order slips: "Annual Report 2026" read 17th, the footers before the figure's caption	photo	the order's chain on a page whose lines, after levelling, sit slightly off their rows	in training, levelled photographs
a footnote mark read as a plain 1	all	plain text has no superscript, and `rec-latin` reads the mark as 1	a later change: mark raised characters, and write them as `<sup>`

Seven of the eleven can be fixed in the engine, with no new training: the block splits, dropping rule lines, trimming the cells, finding the captions and joining the lists. Three are training: the order of tables and of photographed pages, and pictures on photographs. The last, the superscript, is a later change to the recognizer and the writer. The words

themselves are almost all right; what is left is the page's structure.

09 Three pages, and 2,937

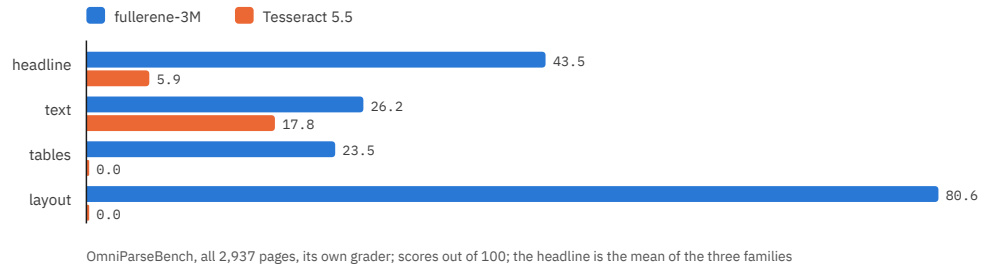


Figure 5. OmniParseBench, all 2,937 pages: fullerene-3M scores a headline of 43.5 against Tesseract 5.5's 5.9; text 26.2 against 17.8, tables 23.5 and layout 80.6 against 0.

Three invented pages show what the engine does; they do not measure how well. That is OmniParseBench, read end to end on all its 2,937 pages and graded by its own grader: `fullerene-3M` scores a headline of 43.5 against Tesseract 5.5's 5.9 (text 26.2 against 17.8; tables 23.5 and layout 80.6, which Tesseract, writing text alone, scores 0), and reads 4.06 pages a second on one core against 0.76 (the model card has every score). Tables are where it leaves the most points, which is what these pages show: tables found and their cells placed, but with rules in their cells, captions as rows, and their place in the page's order.

10 Run it yourself

```
# the set, from Hugging Face
pip install huggingface_hub
huggingface-cli download CarbonResearchLabs/fullerene-3M --local-dir models/fullerene-3M

# the engine's Python module (the README's Build section), then numpy and Pillow
python docs/demo/run.py models/fullerene-3M    # the pages, the output, the numbers
python docs/demo/render.py                      # the pictures, dark and light

# or one page from the command line
carbon-ocr ocr docs/demo/out/page.png --set models/fullerene-3M --format html
```

`run.py` prints the pages with headless Chrome (set `CHROME` if it is not at the Windows default), makes the photo with a fixed seed, runs the engine on one worker and on six, and writes `page.json`, `photo.json` and `tables.json` (the engine's pages: every line's box, text and confidence, the blocks, the cells, the text and the HTML) and `metrics.json` (the numbers above). Those four files, as this run wrote them, are in demo/out/. `render.py` draws the pictures from them, in the Carbon theme, at their size into `docs/assets/` and, with `--big DIR`, at twice it. On another machine the words and the structure come out the same; the times follow its cores and clock.

THIS PAGE IS INTENTIONALLY BLANK

Carbon OCR

Every line read, every page in order.

Every number measured, in public.

EVERY PICTURE DRAWN

From the engine's own output for the page: its lines, its blocks and their order, its cells and its HTML, nothing placed by hand.

EVERY FAULT LISTED

With the stage it comes from and what fixes it, in the engine or in training.

EVERY RUN REPEATABLE

The pages, the photo's seed and the scripts are in the repository; one command makes the numbers and the pictures again.

What the pages show

2.77 M parameters on a laptop's CPU: every line found, 253 of 255 words right, every cell placed.

What is left is the page's structure, and most of it is fixed in the engine.

Karan Vora, Founder · Carbon Research

Carbon-OCR is the engine and **Fullerene** its models. This document lives in the repository as `docs/DEMO.md`, beside the scripts in `docs/demo/` that made its pictures and numbers.

fullerene-3M from huggingface.co/CarbonResearchLabs/fullerene-3M, under the Carbon-OCR OpenRAIL-AMS license, read by the engine of 3 November 2026 on an Intel Core i7-13620H laptop, CPU only. The engine: github.com/karanvora2599/Carbon-OCR.

ORIENT DETECT READ ORDER WRITE

Karan Vora, Founder · Carbon Research